

Eddy Ndumia

AI enthusiast & data engineer (AWS) - Mathematics & Computer Science, JKUAT - Nairobi, Kenya

eddyndumia.github.io/Eddy-Ndumia-Personal-Portfolio | github.com/eddyndumia | linkedin.com/in/eddyndumia

SUMMARY

AI enthusiast with data engineering experience on AWS, turning raw, messy data into pipelines and models that hold up in production. Currently building PesaScore, a fintech platform that scores creditworthiness from alternative data (M-Pesa statements) instead of a traditional credit bureau, and Capstone, an AI lead assistant for real estate agents. Freelance web developer since 2020.

EXPERIENCE

Independent Builder - PesaScore, Capstone & Chukua

2026 - Present

- PesaScore: alt-data credit scoring from M-Pesa statements. Parser tested on a real 7,400-transaction statement, explainable 300-850 scorecard, borrower consent, Postgres RLS. Flutter app, FastAPI, lender web portal.
- Capstone: AI agent handling WhatsApp/Instagram DMs for real estate agencies - qualifies leads, books viewings, hands off to a human when it should. FastAPI + Postgres backend, Next.js dashboard.
- Chukua: Flutter app for giving away free stuff nearby, with AI value estimates and M-Pesa claim fees. Supabase + PostGIS, atomic claim holds, Row-Level Security keeping givers' phones private.

Freelance Web Developer - Kenya

2020 - Present

- Famyard Enterprises (land sales, Nyeri/Laikipia): rebuilt their site in Next.js on Cloudflare Pages, plus a custom CMS (Postgres, role-based login, inquiries inbox) so staff manage plots and content without WordPress.
- Sites and full-stack builds for other small businesses, in WordPress or code.

EDUCATION & TRAINING

B.Sc. Mathematics and Computer Science - JKUAT

2020 - 2026 (graduating December 2026)

ALX / Holberton Software Engineering Program

2023 - low-level programming in C, algorithms, Linux/DevOps

SELECTED PROJECTS

Sales ETL Pipeline - Messy raw sales data cleaned, tested, and loaded into a quality-gated SQLite warehouse; 6 automated data-quality checks gate every run.

Sales Demand Forecasting - Seasonal-naive vs. SARIMA vs. XGBoost on a 90-day holdout; XGBoost's calendar features catch the yearly retail spikes SARIMA misses.

Full write-ups and repo links: eddyndumia.github.io/Eddy-Ndumia-Personal-Portfolio/projects

SKILLS

Languages: Python, TypeScript, JavaScript/Node.js, Dart, SQL

Data & ML: Pandas, NumPy, Scikit-learn, XGBoost, statsmodels, LLM agents (Claude API)

Infra: Postgres, PostGIS, Supabase, Docker, AWS

Web & mobile: React, Next.js, Flutter, FastAPI, WordPress